

## ACM Data Camp 2014

# Principal Component Analysis with Applications in R and Python

*Irina Kukuyeva, Ph.D.*  
ikukuyeva@myvps.org

October 25, 2014



# Outline

- 1 Introduction
- 2 Methodology
- 3 Applications of PCA
- 4 Extensions of PCA
- 5 Final Remarks

- 1 Introduction
  - Prerequisites
  - Motivation
- 2 Methodology
- 3 Applications of PCA
- 4 Extensions of PCA
- 5 Final Remarks

## Course Prerequisite

Introductory class in statistics or linear algebra.

## Software Prerequisites

- Working knowledge of R and/or Python.
- Completed installation of either set of packages:

R

```
library(lattice)  
library(blockcluster) # optional
```

Python

```
import matplotlib.pyplot as plt  
import numpy as np  
import os  
import pandas as pd  
import scipy  
from sklearn.decomposition import PCA
```



## Why use a Dimension Reduction Method?


- I have too many variables... (to model, visualize, etc.).
- Some of my variables are (really) noisy.
- A few of my variables are redundant.
- I want to uncover patterns in the data.

## When to use a Dimension Reduction Method?

- Standalone analysis
- Preprocessing for:
  - Regression
  - Clustering
  - etc.

# Types of Dimension Reduction Methods

- ① Directly observed data-generating process(es).
  - Reduce dimensionality to remove noisy and/or multicollinear variables.
- ② Indirectly observed data-generating process(es).
  - Reduce dimensionality to uncover latent variables and reduce noise.



## 1 Introduction

## 2 Methodology

- Assumptions
- Derivation of PCA
- Determining the Number of Components to Extract
- Residual Analysis
- Interpreting the Components

## 3 Applications of PCA

## 4 Extensions of PCA

## 5 Final Remarks



## Assumption of PCA

- ➊ Observed variables are linear combinations of unobserved *uncorrelated* variables (called *Principal Components*).
- ➋ Variable means and the variance-covariance matrix capture all of the information in the data set.
- ➌ The variances of variables are similar<sup>a</sup>.
  - e.g. Perform PCA on the correlation matrix.
- ➍ Principal Components are unique up to sign and permutation [4].

---

<sup>a</sup>If the variances of the variables are not similar, transform the variables; otherwise those variables with largest variances will seem most important.

## Notation for PCA

- Let  $\mathbf{X}_{N \times M} = \begin{pmatrix} x_{11} & \dots & x_{1M} \\ \vdots & \ddots & \vdots \\ x_{N1} & \dots & x_{NM} \end{pmatrix} = (\mathbf{X}_1 \dots \mathbf{X}_M)$  be the data matrix.
- Let  $\mathbf{\Sigma}$  be the full rank, sample variance-covariance matrix of  $\mathbf{X}$  [4]<sup>1</sup>.
- Let  $\mathbf{Y}$  denote the principal components:

$$Y_1 = \mathbf{X}V_1 = v_{11}X_1 + v_{21}X_2 + \dots + v_{M1}X_M$$

$$Y_2 = \mathbf{X}V_2 = v_{12}X_1 + v_{22}X_2 + \dots + v_{M2}X_M$$

•  
•  
•

$$Y_M = \mathbf{X}V_M = v_{1M}X_1 + v_{2M}X_2 + \dots + v_{MM}X_M$$

<sup>1</sup>Please see [4], Section 3.4 for a discussion on data sets with zero and/or equal variances.

# Notation for PCA

- In matrix form, the principal components  $\mathbf{Y}$  are defined as:

$$\mathbf{Y} = \mathbf{XV} \quad (1)$$

- By assumption (3) of PCA,

$$\text{Var}(Y_i) = V_i^T \Sigma V_i, \quad i = 1, 2, \dots, M \quad (2)$$

$$\text{Cov}(Y_i, Y_j) = V_i^T \Sigma V_k = 0, \quad i, k = 1, 2, \dots, M, \quad i \neq k \quad (3)$$

# Intuition Behind the Derivation I

## Goal of PCA

Maximize the variance explained by each of the principal components<sup>a</sup>:

$$\max_{V_i} \text{Var}(Y_i) = \max_{V_i} V_i^T \Sigma V_i$$

**Solution 1** Set  $V_i = \infty$ .

**Solution 2** Add a constrain to the length of  $V_i$  and solve:


$$\max_{V_i} \text{Var}(Y_i) \text{ subject to } V_i^T V_i = 1$$

---

<sup>a</sup>By assumptions of PCA,  $\text{Cov}(Y_i, Y_j) = 0$ ,  $i \neq j$ .

# Intuition Behind the Derivation II

## Visualization of PCA [1]





## Derivation

**Step 3.** Rearrange the terms of Equation 5 [4]:

$$\begin{aligned} (\mathbf{\Sigma} - \lambda_1 \mathbf{I}_M) \mathbf{V}_1 &= 0 \\ &\vdots \\ (\mathbf{\Sigma} - \lambda_M \mathbf{I}_M) \mathbf{V}_M &= 0, \end{aligned} \quad (6)$$

where  $I_M$  is the  $M \times M$  identity matrix.

# Derivation

**Step 4.** Find the solution to the maximization problem [4].

We know that  $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_M \geq 0$ .

- 1 A one-component solution is *attained* for  $\lambda_1$  (the largest eigenvalue), with the corresponding eigenvector  $V_1$ .

$\vdots$

- 2 An  $L$ -component solution is *attained* for  $\lambda_1, \dots, \lambda_L$  (for  $L, L < M$ , largest eigenvalues), with the corresponding eigenvectors  $V_1, \dots, V_L$ .





# How Many Components to Extract?

Most commonly-employed *ad-hoc* rules

- *A priori*, determine a (cumulative) percent of total variance in the data set that  $L$  components should explain [4].
  - Usually between 70% and 90%.
- Look for an elbow in a scree plot; its occurrence is the number of components we should use [4].
- Keep components that have a contribution of  $> 5\%$ .
- Retain components that explain a 
$$> \frac{1}{\text{number of variables in the data}} \times 100\%$$
 of the variance in the original data set.
- Employ cross-validation to calculate the PREdiction Sum of Squares (PRESS) statistic [4].

# Evaluating Model Fit: Residual Analysis

Suppose we keep  $L$  ( $L < M$ ) components, then

- $\mathbf{V}_0$  is an  $M \times L$  matrix with  $L$  eigenvectors corresponding to the  $L$  largest eigenvalues;
- $\mathbf{Y}_0$  is an  $N \times L$  matrix with  $L$  principal components;
- residuals are:

$$\begin{aligned}\boldsymbol{\Psi} &= \mathbf{X} - \hat{\mathbf{X}} \\ &= \mathbf{X} - \mathbf{Y}_0 \mathbf{V}_0^T \\ &= \mathbf{X} - (\mathbf{X} \mathbf{V}_0) \mathbf{V}_0^T\end{aligned}\tag{7}$$

# Interpreting the Components

- Analyze the weights of the raw variables  $V_i$  used to form principal components to explain the variation in the data:
  - Which variables carry the most weight?
  - Which variables are equally weighted?
  - Which variables offset others (i.e. equally weighted but of opposite signs)?

## 1 Introduction

## 2 Methodology

## 3 Applications of PCA

- Example 1: Food consumption in European Countries
- Example 2: Pasadena Meetup.com Groups

## 4 Extensions of PCA

## 5 Final Remarks

## Data Set Overview

- 1970s food consumption of 25 European countries [11], [8]
- Consumption categorized into 9 food groups [8]:
  - red meat
  - white meat
  - eggs
  - milk
  - fish
  - cereal
  - starch
  - nuts
  - fruits and vegetables
- Units are in 0.1 grams of protein, per person, per day [10].

## Example 1: Food consumption in European Countries

## Example 1: European Food Consumption I

## Step 1: Read the Dataset into R

```
url = 'http://www.stat.ucla.edu/~rgould/120bs06/protein.txt'
```

## R

```
data.orig <- read.table(url,
                        header=TRUE,
                        sep="\t")

head(data.orig)

# For PCA analysis, keep all the variables
# except the first column with country names:
data <- data.orig[, -1]
```

## Python

```
data_orig = pd.read_csv(url,
                        sep='\t')

data_orig.head()

# For PCA analysis, keep all the variables
# except the first column with country names:
columns_keep = data_orig.columns.values.tolist()[1:]
data = data_orig[columns_keep]
```

## Example 1: Food consumption in European Countries

## Example 1: European Food Consumption II

## Step 1: Read the Dataset into R

|   | Country        | RedMeat | WhiteMeat | Eggs | Milk | Fish | Cereals | Starch | Nuts | Fr.Veg |
|---|----------------|---------|-----------|------|------|------|---------|--------|------|--------|
| 1 | Albania        | 10.1    | 1.4       | 0.5  | 8.9  | 0.2  | 42.3    | 0.6    | 5.5  | 1.7    |
| 2 | Austria        | 8.9     | 14.0      | 4.3  | 19.9 | 2.1  | 28.0    | 3.6    | 1.3  | 4.3    |
| 3 | Belgium        | 13.5    | 9.3       | 4.1  | 17.5 | 4.5  | 26.6    | 5.7    | 2.1  | 4.0    |
| 4 | Bulgaria       | 7.8     | 6.0       | 1.6  | 8.3  | 1.2  | 56.7    | 1.1    | 3.7  | 4.2    |
| 5 | Czechoslovakia | 9.7     | 11.4      | 2.8  | 12.5 | 2.0  | 34.3    | 5.0    | 1.1  | 4.0    |
| 6 | Denmark        | 10.6    | 10.8      | 3.7  | 25.0 | 9.9  | 21.9    | 4.8    | 0.7  | 2.4    |

Figure : First 6 rows of the 'Protein' data set.



## Example 1: Food consumption in European Countries

## Example 1: European Food Consumption I

## Step 2a: Perform Preliminary Diagnostics Graphically

R

```
par(las=2,
    cex.axis=0.8,
    mfrow=c(1,2),
    mai=c(1,0.5,1,0.1)
)
boxplot(data,
         main="'Protein' data set.",
         col=topo.colors(ncol(data))
)
boxplot(data.frame(cor(data)),
         main="cor('Protein') data set.",
         col=topo.colors(ncol(data))
)
```

Python

```
def boxplot(dataset, ax, colors):
    box = ax.boxplot(dataset,
                      patch_artist=True)
    for bx, color in zip(box['boxes'], colors):
        bx.set_facecolor(color)
    for median in box['medians']:
        median.set_linewidth(2)
    ax.set_xticklabels(data.columns.values.tolist(),
                       rotation=90,
                       size=8)

colors = ["#4c00ff", "#004cff", "#00e5ff",
          "#00ff4d", "#4dff00", "#e6ff00",
          "#ffff00", "#ffde59", "#ffe0b3"]


fig = plt.figure(1, figsize=(12,6))
ax1 = fig.add_subplot(121)
boxplot(data.values, ax1, colors)
ax2 = fig.add_subplot(122)
boxplot(data.corr().values, ax2, colors)
fig.show()
```

## Example 1: Food consumption in European Countries


## Example 1: European Food Consumption II

## Step 2a: Perform Preliminary Diagnostics Graphically

'Protein' data set.



cor('Protein') data set.



## Example 1: Food consumption in European Countries

# Example 1: European Food Consumption

## Step 2b: Perform Preliminary Diagnostics (more) Rigorously

R

```
bartlett.test(data.frame(cor(data)))

##
## Bartlett test of homogeneity of variances
##
## data:  data.frame(cor(data))
## Bartlett's K-squared = 2.879, df = 8, p-value = 0.9417
```

Python

```
tmp = np.array(data.corr())
scipy.stats.bartlett(tmp[0], tmp[1], tmp[2],
                    tmp[3], tmp[4], tmp[5],
                    tmp[6], tmp[7], tmp[8])

## (2.878558924474143, 0.94174577322949315)
```

## Example 1: Food consumption in European Countries

# Example 1: European Food Consumption

## Step 3a: Perform PCA

### R

```
out.cor      <- princomp(data, cor=TRUE)

lambda_perc <- out.cor$sdev^2/sum(out.cor$sdev^2)
V           <- out.cor$loadings
Y           <- cor(data) %*% V      # Y = XV
PC          <- out.cor$scores      # PC = (standardized dataset)V
### Compare with ?prcomp and ?svd
```

### Python

```
data_std     = data/data.std()

out_cor      = PCA().fit(data_std)

lambda_perc  = out_cor.explained_variance_ratio_
V            = pd.DataFrame(out_cor.components_.T)
Y            = pd.DataFrame(np.dot(data.corr(), V)) # Y = XV
PC           = pd.DataFrame(out_cor.fit_transform(data_std))
### Compare with 'scipy.linalg',
###              'numpy.linalg' and
###              'matplotlib.mlab.PCA'
```

# Example 1: Food consumption in European Countries

## Example 1: European Food Consumption

### Step 3b: Check the Math of PCA

|           | RedMeat     | WhiteMeat  | Eggs        | Milk       | Fish        | Cereals     | Starch      | Nuts        | Fr.Veg      |
|-----------|-------------|------------|-------------|------------|-------------|-------------|-------------|-------------|-------------|
| RedMeat   | 1.00000000  | 0.15300271 | 0.58560895  | 0.50293111 | 0.06095745  | -0.49987746 | 0.13542594  | -0.34944866 | -0.07422123 |
| WhiteMeat | 0.15300271  | 1.00000000 | 0.62040916  | 0.2814839  | -0.23400923 | -0.41379691 | 0.31377205  | -0.6349618  | -0.06131670 |
| Eggs      | 0.58560895  | 0.6204092  | 1.00000000  | 0.5755331  | 0.06557136  | -0.71243682 | 0.45223071  | -0.5597810  | -0.04551755 |
| Milk      | 0.50293110  | 0.2814839  | 0.57553312  | 1.00000000 | 0.13788370  | -0.59273662 | 0.22241118  | -0.6210875  | -0.40836414 |
| Fish      | 0.06095745  | -0.2340092 | 0.06557136  | 0.1378837  | 1.00000000  | -0.52423080 | 0.40385286  | -0.1471529  | 0.26613865  |
| Cereals   | -0.49987746 | -0.4137969 | -0.71243682 | -0.5927366 | -0.52423080 | 1.00000000  | -0.53326231 | 0.6509973   | 0.04654808  |
| Starch    | 0.13542594  | 0.3137721  | 0.45223071  | 0.2224112  | 0.40385286  | -0.53326231 | 1.00000000  | -0.4743116  | 0.08440956  |
| Nuts      | -0.34944855 | -0.6349618 | -0.55978097 | -0.6210875 | -0.14715294 | 0.65099727  | -0.47431155 | 1.00000000  | 0.37496971  |
| Fr.Veg    | -0.07422123 | -0.0613167 | -0.04551755 | -0.4083641 | 0.26613865  | 0.04654808  | 0.08440956  | 0.3749697   | 1.00000000  |

Figure : Correlation Matrix of the 'Protein' data set.

|           | Comp.1     | Comp.2      |
|-----------|------------|-------------|
| RedMeat   | -0.3026094 | 0.05625165  |
| WhiteMeat | -0.3105562 | 0.23685334  |
| Eggs      | -0.4266785 | 0.03533576  |
| Milk      | -0.3777273 | 0.18458877  |
| Fish      | -0.1356499 | -0.64681970 |
| Cereals   | 0.4377434  | 0.23348508  |
| Starch    | -0.2972477 | -0.35282564 |
| Nuts      | 0.4203344  | -0.14331056 |
| Fr.Veg    | 0.1104199  | -0.53619004 |

Figure : First two loadings.

|           | Comp.1     | Comp.2      |
|-----------|------------|-------------|
| RedMeat   | -1.2123856 | 0.09197142  |
| WhiteMeat | -1.2442241 | 0.38725507  |
| Eggs      | -1.7094608 | 0.05777394  |
| Milk      | -1.5133408 | 0.30180253  |
| Fish      | -0.5434728 | -1.05754986 |
| Cereals   | 1.7537917  | 0.38174798  |
| Starch    | -1.1909042 | -0.57686972 |
| Nuts      | 1.6840435  | -0.23431269 |
| Fr.Veg    | 0.4423904  | -0.87667042 |

Figure : First two Principal Components.

## Example 1: Food consumption in European Countries

## Example 1: European Food Consumption I

## Step 4: Determine How Many Components to Keep

## R

```
par( mai=c(1,1,0.1,0.1) )
plot(lambda_perc,
     type="b",
     xlab="Component Number",
     ylab="Percent total variance explained",
     pch=16
)
text(x=1:ncol(cor(data)),
     y=lambda_perc,
     labels=round(cumsum(lambda_perc), 2),
     adj=-0.3
)
```

## Python

```
plt.plot(range(out_cor.n_components_),
         lambda_perc,
         'o-'
        )
plt.xlabel("Component Number")
plt.ylabel("Percent total variance explained")


for i in range(out_cor.n_components_):
    plt.text(x=i,
             y=lambda_perc[i],
             s=str(round(lambda_perc[i], 2))
            )

plt.show()
```

## Example 1: Food consumption in European Countries

## Example 1: European Food Consumption II

## Step 4: Determine How Many Components to Keep




## Example 1: Food consumption in European Countries

## Example 1: European Food Consumption I

## Step 5: Interpret the Results

R

```
biplot(out.cor,
       pc.biplot=TRUE,
       xlab=data.orig$Country,
       xlim=c(-3, 3))
abline(h=0,
       lty=2,
       col="grey50")
abline(v=0,
       lty=2,
       col="grey50")
```





## Example 1: Food consumption in European Countries

# Example 1: European Food Consumption II

## Step 5: Interpret the Results

### Python

```
plt.plot(PC[[0]],  
         -PC[[1]],  
         'ko'  
        )  
plt.xlabel("PC1")  
plt.ylabel("-PC2")  
plt.axvline(x=0, color='0.75')  
plt.axhline(y=0, color='0.75')  
  
# Add components' labels:  
for i in range(len(PC)):  
    plt.text(x=PC[0][i],  
            y=-PC[1][i],  
            s=data_orig['Country'][i]  
            )
```

```
# Add loading directions:  
for i in range(out_cor.n_components_):  
    plt.arrow(x=0,  
              y=0,  
              dx=1.5*V[0][i],  
              dy=-1.5*V[1][i],  
              color='r',  
              head_width=0.1  
             )  
  
# Add loading labels:  
col_names = data.columns.values.tolist()  
for i in range(out_cor.n_components_):  
    plt.text(x=2*V[0][i],  
            y=-2*V[1][i],  
            s=col_names[i],  
            color='r'  
            )  
  
plt.show()
```

## Example 1: Food consumption in European Countries

## Example 1: European Food Consumption

## Step 5b: Check the Math for Portugal

## ① Original values:

| ##    | RedMeat | WhiteMeat | Eggs | Milk | Fish | Cereals | Starch | Nuts | Fr.Veg |
|-------|---------|-----------|------|------|------|---------|--------|------|--------|
| ## 17 | 6.2     | 3.7       | 1.1  | 4.9  | 14.2 | 27      | 5.9    | 4.7  | 7.9    |

## ② Standardized values:

| ## | RedMeat | WhiteMeat | Eggs    | Milk    | Fish   | Cereals | Starch |
|----|---------|-----------|---------|---------|--------|---------|--------|
| ## | -1.0839 | -1.1359   | -1.6428 | -1.7187 | 2.9143 | -0.4782 | 0.9938 |
| ## | Nuts    | Fr.Veg    |         |         |        |         |        |
| ## | 0.8199  | 2.0866    |         |         |        |         |        |

## ③ Weights of variables forming first component:

| ## | RedMeat | WhiteMeat | Eggs    | Milk    | Fish    | Cereals | Starch  |
|----|---------|-----------|---------|---------|---------|---------|---------|
| ## | -0.3026 | -0.3106   | -0.4267 | -0.3777 | -0.1356 | 0.4377  | -0.2972 |
| ## | Nuts    | Fr.Veg    |         |         |         |         |         |
| ## | 0.4203  | 0.1104    |         |         |         |         |         |

## ④ Multiply (2) by (3) and compare with Principal Component 1 (1.741) for Portugal.

## Example 1: Food consumption in European Countries

## Example 1: European Food Consumption I

## Step 6: Analyze Residuals

$$\text{Residuals} = \mathbf{X} - (\mathbf{XV}_0)\mathbf{V}_0^T$$

R

```
X_hat <- cor(data) -
  (cor(data) %*% V[, 1:2]) %*% t(V[, 1:2])

n_obs <- nrow(X_hat) * ncol(X_hat)

levelplot(X_hat,
  col.regions=heat.colors( n_obs ),
  scales=list( x=list(rot = 90) )
)
```

Python

```
X_hat = data.corr() -
  np.dot(
    np.dot(
      data.corr(),
      pd.DataFrame([V[0], -V[1]]).T
    ),
    pd.DataFrame([V[0], -V[1]])
  )


fig = plt.figure(1, figsize=(12,6))
ax = fig.add_subplot(111)
heatmap = plt.imshow(X_hat, interpolation='none')
heatmap.set_cmap('hot')
ax.set_xticks(range(out_cor.n_components_))
ax.set_yticks(range(out_cor.n_components_))
ax.set_xticklabels(data.columns.values.tolist(),
  rotation=90,
  size=10)
ax.set_yticklabels(data.columns.values.tolist(),
  size=10)

plt.colorbar()
plt.show()
```

## Example 1: Food consumption in European Countries

## Example 1: European Food Consumption II

## Step 6: Analyze Residuals



## Example 2: Pasadena Meetup.com Groups

# Example 2: Pasadena Meetup.com Groups I

## Data Set Creation

1. Scrape 200 groups from Meetup.com, meeting within 25 miles of Pasadena (in Sept 2014)
2. Extract keywords from each group
  - remove quotes and dashes, convert words to lower case
3. Perform word stemming on the key words from each group
  - 'network' and 'networking' are mapped to one word
4. Create a matrix of word counts of keywords from each group (with  $K$  unique words)
5. Compute the distance between each Meetup.com group:

$$\text{Distance b/w meetups } i \text{ and } j = \frac{\sum_{k=1}^K |\text{count}_{k,i} - \text{count}_{k,j}|}{K}$$

## Example 2: Pasadena Meetup.com Groups

## Example 2: Pasadena Meetup.com Groups II

## Data Set Creation

## Processing keywords of one meetup group:

## We're about:

Small Business ·  
 Self-Improvement · Social  
 · Fun Times · New In  
 Town · Dating and  
 Relationships ·  
 Professional Networking ·  
 Singles · Speed Dating ·  
 Business Referral  
 Networking · Friendship  
 and Connections · Small  
 Business Networking ·  
 Entrepreneur Networking  
 · Social Networking ·  
 Women's Business  
 Networking

| word counts | words        |
|-------------|--------------|
| 6           | network      |
| 4           | busi         |
| 2           | and          |
| 2           | date         |
| 2           | small        |
| 2           | social       |
| 1           | connect      |
| 1           | entrepreneur |
| 1           | friendship   |
| 1           | fun          |
| 1           | in           |
| 1           | new          |
| 1           | profession   |
| 1           | referr       |
| 1           | relationship |
| 1           | selfimprov   |
| 1           | singl        |
| 1           | speed        |
| 1           | time         |
| 1           | town         |
| 1           | women        |

## Example 2: Pasadena Meetup.com Groups

# Example 2: Pasadena Meetup.com Groups

## Step 1: Read the Dataset into R

R

```
dist_mat <- read.csv("../Data/meetup_dist_mat.csv",  
                      row.names=1  
                      )
```

Python

```
os.chdir("../Data")  
dist_mat = pd.read_csv('meetup_dist_mat.csv',  
                       sep=',',  
                       index_col=0  
                       )
```

## Example 2: Pasadena Meetup.com Groups

# Example 2: Pasadena Meetup.com Groups

## Step 2: Perform PCA

### R

```
out.cor      <- princomp(dist_mat, cor=TRUE)

lambda_perc  <- out.cor$sdev^2/sum(out.cor$sdev^2)
V            <- out.cor$loadings
Y            <- cor(dist_mat) %*% V      # Y = XV
PC           <- out.cor$scores          # PC = (standardized dataset)V
```

### Python

```
data_std = dist_mat/dist_mat.std()
out_cor = PCA()      # keep all components
out_cor.fit(data_std) # fit PCA

lambda_perc = out_cor.explained_variance_ratio_
V           = pd.DataFrame(out_cor.components_.T)
Y           = pd.DataFrame(np.dot(dist_mat.corr(), V))
PC          = pd.DataFrame(out_cor.fit_transform(data_std))
```



# Example 2: Pasadena Meetup.com Groups I

## Step 3: Determine How Many Components to Keep

### R

```
lambda_perc_cumsum <- cumsum(lambda_perc)

plot(lambda_perc[2:20],
     type="b",
     xlab="Component Number (2-20)",
     ylab="% total variance explained",
     axes=FALSE)
axis(1,
     at=1:20,
     lab=2:21,
     las=2)
box()
text(x=2:20,
     y=lambda_perc[2:20],
     labels=round(
         lambda_perc_cumsum[2:20],
         digits=2
     ),
     adj=-0.1,
     pos=3,
     cex=0.7
)
```

### Python

```
lambda_perc_cumsum = np.cumsum(lambda_perc)

fig = plt.figure()
ax = fig.add_subplot(111)

plt.plot(range(1, 21),
         lambda_perc[1:21], 'ko-')
plt.xticks(np.array(range(1, 21)))
ax.set_xticklabels(range(1, 21), rotation=90)
ax.set_yticklabels([])
ax.set_xlabel("Component Number (2 through 20)")
ax.set_ylabel("Percent total variance explained")


for i in range(1, 21):
    ax.text(x=i-0.2,
           y=lambda_perc[i],
           s=str(round(lambda_perc_cumsum[i], 2)),
           size='small'
          )

fig.show()
```

## Example 2: Pasadena Meetup.com Groups

## Example 2: Pasadena Meetup.com Groups II

## Step 3: Determine How Many Components to Keep



# Example 2: Pasadena Meetup.com Groups I

## Step 4: Interpret the Results

### R

```
### --- Look at relationships between weights
V_first4 <- V[, 1:4]
### --- Look at relationships between components
Y_first4 <- Y[, 1:4]
```

### Python

```
### --- Look at relationships between weights
V_first4 = pd.DataFrame([-V[0], V[1], -V[2], -V[3]]).T

### --- Look at relationships between components
Y_first4 = pd.DataFrame([-Y[0], Y[1], -Y[2], -Y[3]]).T
```

Please see 'First\_four\_loadings.PDF' for a table of Meetup groups with  $|loadings| \geq 0.1$ .

## Example 2: Pasadena Meetup.com Groups

# Example 2: Pasadena Meetup.com Groups II

## Step 4: Interpret the Results

First four loadings can be interpreted as follows:

Component 1 None

Component 2 social/networking (pos weights) versus  
fitness/yoga (neg weights)

Component 3 personal development (pos weights) versus  
kids and pets (neg weights)

Component 4 networking/pets (pos weights) versus  
moms (neg weights)

## Example 2: Pasadena Meetup.com Groups

## Example 2: Pasadena Meetup.com Groups I

## Step 5: Alternative Analysis via Co-Clustering

## R

```
# 'nbcocluster' = 4 clusters for rows and 4 for columns
# 'model' = equal-sized clusters with unequal variance
out <- cocluster(cor(dist_mat),
                 datatype="continuous",
                 nbcocluster=c(4,4),
                 model="pi_rho_sigma2kl"
                 )

## Co-Clustering successfully terminated!

plot(out)
```

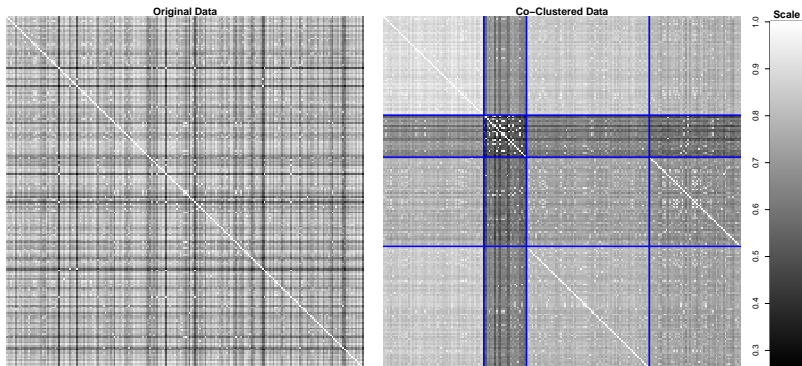
## Python

```
# No packages/functions exist for
# block clustering with Python
```

## Example 2: Pasadena Meetup.com Groups

## Example 2: Pasadena Meetup.com Groups II

## Step 5: Alternative Analysis via Co-Clustering



- 1 Introduction
- 2 Methodology
- 3 Applications of PCA
- 4 Extensions of PCA**
  - Assumptions of PCA
  - Standalone Analysis
  - Cluster Analysis
- 5 Final Remarks

## Recall: Assumptions of PCA

- ❶ *Principal Components* are uncorrelated with each other.
- ❷ Observed variables are combinations of *Principal Components*).
  - a Observed variables are unconstrained.
  - b Observed variables are continuous.
  - c Observed variables are linear combinations of *Principal Components*).
- ❸ Observed variable means and the variance-covariance matrix capture all of the information in the data set.
- ❹ Observed variables are best represented by an  $N \times M$  matrix.
- ❺ *Principal Components* are unique up to sign and permutation.




# Relax Assumption (1)

## Factor Analysis

### Overview: Factor Analysis

Rotates Principal Components by an angle to derive more interpretable factors.

### Visualization of method



### Implementation

**R** `varimax`

**Python** `sklearn.decomposition.FactorAnalysis`

# Relax Assumption (2a)

## Non-negative Matrix Factorization

### Overview: Non-negative Matrix Factorization

Constrains weights/loadings to be positive and finds Principal Components as outlined above.

### Implementation

R    library(NMF)

Python    NIMFA

# Relax Assumption (2b)

## Correspondence Analysis

### Overview: Correspondence Analysis

Application of PCA to discrete variables [4].

### Implementation

R    `library(ca)`

Python    `mca`

# Relax Assumptions 2(b) and (c)

## Nonlinear Principal Component Analysis

### Overview: Nonlinear Principal Component Analysis

Generalize PCA to simultaneously optimize over nonlinear transformations of variables and PCA estimation [2], [7].

### Implementation

**R** `library(homals) library(kernlab),  
library(autoencoder) [9], etc.`

**Python** `help(sklearn.decomposition.KernelPCA), pynnet, etc.`

# Relax Assumption (1) and (3)

## Independent Component Analysis

### Overview: Independent Component Analysis (ICA)

Rotates Principal Components to be independent (i.e. nonlinearly uncorrelated) [3].

- Principal components from multivariate normally distributed data sets are independent.
- Independence can be quantified by:
  - Mutual Information
  - Cumulants – diagonal for independent variables

### Implementation

**R** `library(fastICA), library(PearsonICA)`

**Python** `sklearn.decomposition.FastICA`

# Relax Assumption (4) and partially (5) I

## Tucker- $n$ /CANDECOMP

### Overview: Tucker- $n$ and CANDECOMP

Generalize PCA to data sets with  $n$  modes, satisfying either [5]:

- uniqueness up to permutation via Tucker-3; or
- uniqueness up to scale via CANDECOMP.

### Implementation


R `library(ThreeWay)`

Python `PyTensor`

# Relax Assumption (4) and partially (5) II

## Tucker- $n$ /CANDECOMP

### Visualization of method



# Relax Assumption (1), partially (3), and (4) I

## Array Independent Component Analysis

### Overview: Array Independent Component Analysis [6]

Generalizes of ICA to data sets with  $n$  modes, accounting for relationships between modes.

- Relationships are quantified by cumulants – diagonal for independent variables.


### Implementation

R   Coming soon.



# Relax Assumption (1), partially (3), and (4) II

## Array Independent Component Analysis



# Spectral Clustering

## Overview: Spectral Clustering [12]

- 1 Quantifies similarities between the variables of the original data set in a matrix.
- 2 Reduces the dimensionality of the similarity matrix.
- 3 Clusters the smaller data set, from Step (2).

## Implementation

**R**    `library(kernlab)`

**Python**    `sklearn.cluster.bicluster.SpectralCoclustering`

- 1 Introduction
- 2 Methodology
- 3 Applications of PCA
- 4 Extensions of PCA
- 5 Final Remarks**

# Final Remarks

Python code has been tested with 2.7.5.

R code has been tested with 3.0.3.

Tutorial is available here:

<https://sites.google.com/site/ikukuyeva/>

We are hiring (soon)!

# References I

- [1] cMinor  
(<http://math.stackexchange.com/users/7265/cminor>).  
Visualization of singular value decomposition of a symmetric matrix. (Mathematics), Nov 2012.
- [2] Jan de Leeuw. Nonlinear principal component analysis and related techniques. Technical report, UCLA, 2005.
- [3] A. Hyvärinen, J. Karhunen, and E. Oja. *Independent Component Analysis*. Wiley Series, 2001.
- [4] I. T. Jolliffe. *Principal Component Analysis*. Springer Series in Statistics. Springer, 2002.
- [5] Pieter M. Kroonenberg. *Applied Multiway Data Analysis*. John Wiley and Sons, 2008.

## References II

- [6] Irina Kukuyeva. *Array Independent Component Analysis with Application to Remote Sensing*. PhD thesis, UCLA, 2012.
- [7] M. Linting, J.J. Meulman, P.J.F. Groenen, and A.J. Van der Kooij. *Nonlinear Principal Components Analysis: Nonlinear Principal Component Analysis: Introduction and Application*. American Psychological Association., 2007.
- [8] SAS/STAT(R) 9.22 User's Guide. Creating a distance matrix as input for a subsequent cluster analysis.
- [9] Matthias Scholz. Nonlinear pca.
- [10] StatSoft Inc. Example 3: Protein consumption in europe, 2012.

# References III

- [11] Universidade Federal de Goiás. Datafile Name: Protein.
- [12] Ulrike von Luxburg. A tutorial on spectral clustering.  
*Statistics and Computing*, 17(4), 2007.