

Model checking and model selection

Jeffrey Leek, Assistant Professor of Biostatistics
Johns Hopkins Bloomberg School of Public Health

Model checking and model selection



- Sometimes model checking/selection not allowed
- Often it can lead to problems
 - Overfitting
 - Overtesting
 - Biased inference
- *But* you don't want to miss something obvious

Linear regression - basic assumptions

- Variance is constant
- You are summarizing a linear trend
- You have all the right terms in the model
- There are no big outliers

Model checking - constant variance

```
set.seed(3433); par(mfrow=c(1,2))  
data <- rnorm(100,mean=seq(0,3,length=100),sd=seq(0.1,3,length=100))  
lm1 <- lm(data ~ seq(0,3,length=100))  
plot(seq(0,3,length=100),data,pch=19,col="grey"); abline(lm1,col="red",lwd=3)  
plot(seq(0,3,length=100),lm1$residuals,,pch=19,col="grey"); abline(c(0,0),col="red",lwd=3)
```



What to do

- See if another variable explains the increased variance
- Use the *vcovHC* {sandwich} variance estimators (if n is big)


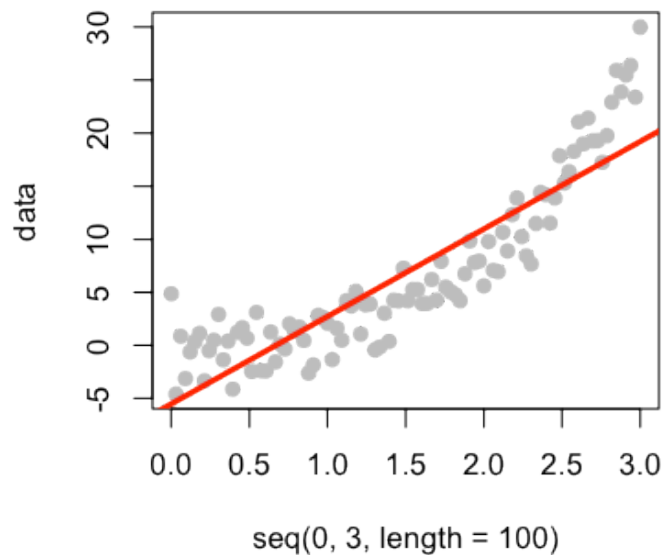
Using the sandwich estimate

```
set.seed(3433); par(mfrow=c(1,2)); data <- rnorm(100,mean=seq(0,3,length=100),sd=seq(0.1,3,length=100))
lm1 <- lm(data ~ seq(0,3,length=100))
vcovHC(lm1)
summary(lm1)$cov.unscaled
```

	(Intercept)	seq(0, 3, length = 100)
(Intercept)	0.03941	-0.01960
seq(0, 3, length = 100)	-0.01960	0.01307

Model checking - linear trend

```
set.seed(3433); par(mfrow=c(1,2))  
data <- rnorm(100,mean=seq(0,3,length=100)^3,sd=2)  
lm1 <- lm(data ~ seq(0,3,length=100))  
plot(seq(0,3,length=100),data,pch=19,col="grey"); abline(lm1,col="red",lwd=3)  
plot(seq(0,3,length=100),lm1$residuals,,pch=19,col="grey"); abline(c(0,0),col="red",lwd=3)
```



What to do

- Use Poisson regression (if it looks exponential/multiplicative)
- Use a data transformation (e.g. take the log)
- Smooth the data/fit a nonlinear trend (next week's lectures)
- Use linear regression anyway
 - Interpret as the linear trend between the variables
 - Use the *vcovHC* {sandwich} variance estimators (if n is big)

Model checking - missing covariate



```
set.seed(3433); par(mfrow=c(1,3)); z <- rep(c(-0.5,0.5),50)
data <- rnorm(100,mean=(seq(0,3,length=100) + z),sd=seq(0.1,3,length=100))
lm1 <- lm(data ~ seq(0,3,length=100))
plot(seq(0,3,length=100),data,pch=19,col=((z>0)+3)); abline(lm1,col="red",lwd=3)
plot(seq(0,3,length=100),lm1$residuals,pch=19,col=((z>0)+3)); abline(c(0,0),col="red",lwd=3)
boxplot(lm1$residuals ~ z,col = ((z>0)+3) )
```

What to do

- Use exploratory analysis to identify other variables to include
- Use the *vcovHC* {sandwich} variance estimators (if n is big)
- Report unexplained patterns in the data

Model checking - outliers

```
set.seed(343); par(mfrow=c(1,2)); betahat <- rep(NA,100)
x <- seq(0,3,length=100); y <- rcauchy(100); lm1 <- lm(y ~ x)
plot(x,y,pch=19,col="blue"); abline(lm1,col="red",lwd=3)
for(i in 1:length(data)){betahat[i] <- lm(y[-i] ~ x[-i])$coeff[2]}
plot(betahat - lm1$coeff[2],col="blue",pch=19); abline(c(0,0),col="red",lwd=3)
```



11/25

What to do

- If outliers are experimental mistakes -remove and document them
- If they are real - consider reporting how sensitive your estimate is to the outliers
- Consider using a robust linear model fit like *rlm*{MASS}

Robust linear modeling

```
set.seed(343); x <- seq(0,3,length=100); y <- rcauchy(100);  
lm1 <- lm(y ~ x); rlm1 <- rlm(y ~ x)  
lm1$coeff
```

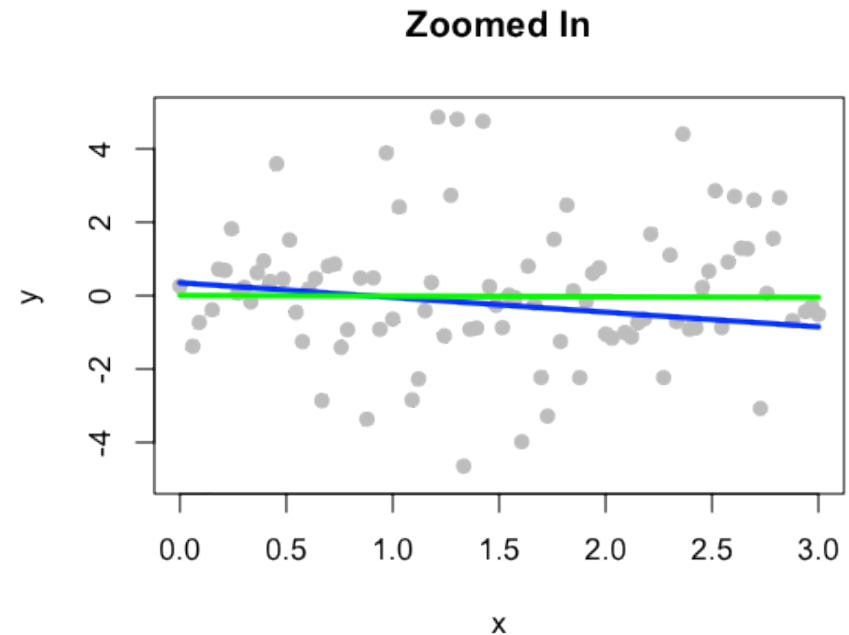

(Intercept)	x
0.3523	-0.4011

```
rlm1$coeff
```

(Intercept)	x
0.008527	-0.017892

Robust linear modeling



```
par(mfrow=c(1,2))
plot(x,y,pch=19,col="grey")
lines(x,lm1$fitted,col="blue",lwd=3); lines(x,rlm1$fitted,col="green",lwd=3)
plot(x,y,pch=19,col="grey",ylim=c(-5,5),main="Zoomed In")
lines(x,lm1$fitted,col="blue",lwd=3); lines(x,rlm1$fitted,col="green",lwd=3)
```



14/25

Model checking - default plots





```
set.seed(343); par(mfrow=c(1,2))  
x <- seq(0,3,length=100); y <- rnorm(100); lm1 <- lm(y ~ x)  
plot(lm1)
```



Model checking - deviance

- Commonly reported for GLM's
- Usually compares the model where every point gets its own parameter to the model you are using
- On it's own it doesn't tell you what is wrong
- In large samples the deviance may be big even for "conservative" models
- You can not compare deviances for models with different sample sizes

R^2 may be a bad summary



Model selection

- Many times you have multiple variables to evaluate
- Options for choosing variables
 - Domain-specific knowledge
 - Exploratory analysis
 - Statistical selection
- There are many statistical selection options
 - Step-wise
 - AIC
 - BIC
 - Modern approaches: Lasso, Ridge-Regression, etc.
- Statistical selection may bias your inference
 - If possible, do selection on a held out sample

Error measures

- R^2 alone isn't enough - more variables = bigger R^2
- [Adjusted \$R^2\$](#) is R^2 taking into account the number of estimated parameters
- [AIC](#) also penalizes models with more parameters
- [BIC](#) does the same, but with a bigger penalty

Movie Data

```
download.file("http://www.rossmanchance.com/iscam2/data/movies03RT.txt",destfile="./data/movies.txt")
movies <- read.table("./data/movies.txt",sep="\t",header=T,quote="")
head(movies)
```

				X score	rating	genre	box.office	running.time
1	2	Fast 2 Furious	48.9	PG-13	action/adventure	127.15	107	
2		28 Days Later	78.2	R	horror	45.06	113	
3		A Guy Thing	39.5	PG-13	rom comedy	15.54	101	
4		A Man Apart	42.9	R	action/adventure	26.25	110	
5		A Mighty Wind	79.9	PG-13	comedy	17.78	91	
6		Agent Cody Banks	57.9	PG	action/adventure	47.81	102	

<http://www.rossmanchance.com/>

Model selection - step

```
movies <- movies[,-1]
lm1 <- lm(score ~ .,data=movies)
aicFormula <- step(lm1)
```

Start: AIC=727.5
 score ~ rating + genre + box.office + running.time

	Df	Sum of Sq	RSS	AIC
- genre	12	2575	22132	721
- rating	3	40	19596	722
- running.time	1	237	19793	727
<none>			19556	728
- box.office	1	3007	22563	746

Step: AIC=720.8
 score ~ rating + box.office + running.time

	Df	Sum of Sq	RSS	AIC
- rating	3	491	22623	718
<none>			22132	721

Model selection - step

```
aicFormula
```

Call:


```
lm(formula = score ~ box.office + running.time, data = movies)
```

Coefficients:

(Intercept)	box.office	running.time
37.2364	0.0824	0.1275

Model selection - regsubsets

```
library(leaps);
regSub <- regsubsets(score ~ .,data=movies)
plot(regSub)
```



<http://cran.r-project.org/web/packages/leaps/leaps.pdf>

Model selection - bic.glm

```
library(BMA)
bicglm1 <- bic.glm(score ~.,data=movies,glm.family="gaussian")
print(bicglm1)
```

Call:

```
bic.glm.formula(f = score ~ ., data = movies, glm.family = "gaussian")
```

Posterior probabilities(%):

rating	genre	box.office	running.time
0.0	100.0	100.0	18.2

Coefficient posterior expected values:

(Intercept)	ratingPG	ratingPG-13	ratingR
45.263	0.000	0.000	0.000
genreaction/adventure	genreanimated	genrecomedy	genredocumentary
-0.120	7.628	2.077	8.642
genredrama	genrefantasy	genrehorror	genremusical
13.041	1.504	-3.458	-12.255

Notes and further resources

- Exploratory/visual analysis is key
- Automatic selection produces an answer - but may bias inference
- You may think about separating the sample into two groups
- The goal is not to get the "causal" model
- [Lars package](#)
- [Elements of machine learning](#)